

DBSCAN

Evan Burdett, Isaac Summerford, Carter Williams


DBSCAN Introduction

- DBSCAN: Density-Based Spatial Clustering of Applications with Noise
- Data clustering algorithm proposed by Ester, Kriegel, Sander, and Xu in their paper published in 1996
- Density Based and Non-Parametric (no specific distribution)
- Given a set of points in some space, closely packed points are grouped together and points in low-density areas are marked as outliers

DBSCAN Introduction: Parameters and Points

- Parameters
 - eps (epsilon): radius of region considered to be a neighborhood
 - If eps is too small, you could end up having too many points be labeled as noise
 - If eps is too big, separate clusters may merge together and defeat the purpose of the algorithm
 - minPts: minimum number of points within the eps radius to make a neighborhood dense
- Points
 - Core Points: points that are within a dense neighborhood
 - Border Points: points that are close to core points but do not have enough neighbors to be considered core points
 - Outlier Points: points that do not belong to any cluster and are considered noise

DBSCAN Introduction: Parameters and Points



$\text{MinPts} = 4$

Red: Core Points

Yellow: Border points. Still part of the cluster because it's within epsilon of a core point, but not does not meet the `min_points` criteria

Blue: Noise point. Not assigned to a cluster


DBSCAN Introduction: Steps

1. Identify Core Points
 - a. Iterate through all of the points, counting the number of points within its eps neighborhood, if it is greater than or equal to minPts it is a core point
2. Form Clusters
 - a. Make a cluster for every core point that does not have one yet, add all points within the radius of the core point to the cluster
3. Establish Density Connectivity
 - a. Two points have density connectivity if there is a chain of points that exists where each point is within the radius of the next point
4. Label Noise Points
 - a. All points that are not in a cluster are labeled as noise

DBSCAN Introduction: Advantages and Disadvantages

- Advantages:
 - Do not need to specify the number of clusters like you have to in K-Means
 - It's good at identifying arbitrarily shaped clusters, other algorithms are only good at identifying spherical clusters
 - Handles noise well when compared to other algorithms
- Disadvantages:
 - Not as good for straightforward and spherical data sets
 - Has two parameters instead of just one in other algorithms

DBSCAN Introduction: DBSCAN vs K-Means



Real-World Examples

- **Geospatial Analysis**
 - Mapping crime incidents in a city to find high-crime “hotspots.”
- **Fraud & Anomaly Detection**
 - Flagging unusual credit card transactions.
- **Customer Segmentation**
 - Grouping shoppers by shopping frequency and spending level
- **Image Processing**
 - Segmenting brain tumors in an MRI scan.
- **Medical & Biological Research**
 - Clustering genes with similar expression levels across many samples.
- **Astronomy**
 - Identifying galaxy clusters in deep-space sky surveys.

Dataset & Model Design


Dataset: Real-world "Mall Customer" data comparing Annual Income to Spending Score (on a 1-100 scale).

Model: DBSCAN (Density-Based Spatial Clustering).

Why this fits our data:


- It automatically labels irregular behavior as "Noise" (outliers).
- It's great for finding natural, non-spherical cluster shapes.
- It discovers the right number of clusters on its own (no need to specify a fixed K).

Core Logic: Points are grouped by density. A cluster starts when a "Core Point" has at least 5 neighbors (MinPts) within a specific distance (Epsilon).




Dataset Overview and Preprocessing

- Dataset: Real-world "Mall Customer" data comparing Annual Income to Spending Score (1–100).
- Preprocessing: Applied StandardScaler to normalize inputs.
 - Income (\$15k–\$137k) and Score (1–100) have vastly different magnitudes.
 - Without scaling, the Euclidean distance calculation would be entirely biased toward Income.
- Computational Efficiency: Used Scikit-Learn's implementation utilizing KD-Trees for spatial indexing.
 - Reduces neighbor search complexity from $O(N^2)$ to $O(N \log N)$.



Parameter Optimization (Finding Epsilon)



- The Elbow Method: Calculated the distance to the 5th nearest neighbor for every datapoint.
 - Sorted and plotted these distances to visualize the density drop-off
 - The “Elbow” (sharp curve) suggested an initial Epsilon of ~ 0.5.
- Iterative Tuning: We manually adjusted epsilon (0.50 -> 0.40 -> 0.365).
 - The goal was to resolve the “Bridge Effect,” where distinct clusters were merged due to density gradients.
 - Result: An epsilon of 0.365 provided the cleanest separation between clusters.



Clustering Results & Evaluation

We evaluated the results using a few key criteria:

- Qualitative Inspection: Since we lacked true labels, we focused on how easily we could interpret the resulting clusters. For example, did shoppers with "High Income/Low Spend" form a clear, distinct group? 5 natural segments.
- Noise Ratio: We monitored the percentage of points labeled as X (Noise).
 - A ratio that was too high (>20%) suggested Epsilon was too strict.
 - A ratio that was too low (<1%) suggested Epsilon was too loose, potentially merging distinct clusters.
 - Our final ratio was 10.50% (21 outliers)
- Cluster Count: We verified if the model successfully identified the 5 natural segments we had observed in the data's distribution.



Our Code

```
# Final DBSCAN run and evaluation

db_final = DBSCAN( variable ) labels_ ndarray[_AnyShape, dtype[Any]]
labels = db_final.labels_

n_clusters = len(set(labels)) - (1 if -1 in labels else 0)
n_noise = list(labels).count(-1)
noise_ratio = n_noise / len(labels)

print("Results:")
print(f"Clusters found: {n_clusters}")
print(f"Noise points (outliers): {n_noise}")
print(f"Noise ratio: {noise_ratio:.2%}")

plt.figure(figsize=(10, 6))


unique_labels = set(labels)
colors = [plt.cm.Spectral(each) for each in np.linspace(0, 1, len(unique_labels))]

for k_label, col in zip(unique_labels, colors):
    if k_label == -1:
        col = (0, 0, 0, 1)
        label_text = "Noise (outliers)"
        marker_shape = "x"
        size = 100
    else:
        label_text = f"Cluster {k_label}"
        marker_shape = "o"
        size = 50


    # nothing special, same col

    class_member_mask = labels == k_label
    xy = X_scaled[class_member_mask]
    plt.scatter(
        xy[:, 0], xy[:, 1], c=[col], marker=marker_shape, s=size, label=label_text
    )

plt.title("DBSCAN clustering Results (Mall Customers)")
plt.xlabel("Annual Income (scaled)")
plt.ylabel("Spending Score (scaled)")
plt.legend(bbox_to_anchor=(1.05, 1), loc="upper left")
plt.grid(True, alpha=0.3)
plt.tight_layout()
plt.show()
```



Demo of Our Code



Sources

GeeksforGeeks. “DBSCAN Clustering in ML Density Based Clustering.” *GeeksforGeeks*, 6 May 2019,

www.geeksforgeeks.org/machine-learning/dbscan-clustering-in-ml-density-based-clustering/.

Lutins, Evan. “DBSCAN: What Is It? When to Use It? How to Use It.” *Medium*, 4 Dec. 2020,

elutins.medium.com/dbscan-what-is-it-when-to-use-it-how-to-use-it-8bd506293818.

Wikipedia Contributors. “DBSCAN.” *Wikipedia*, Wikimedia Foundation, 16 July 2019, en.wikipedia.org/wiki/DBSCAN.